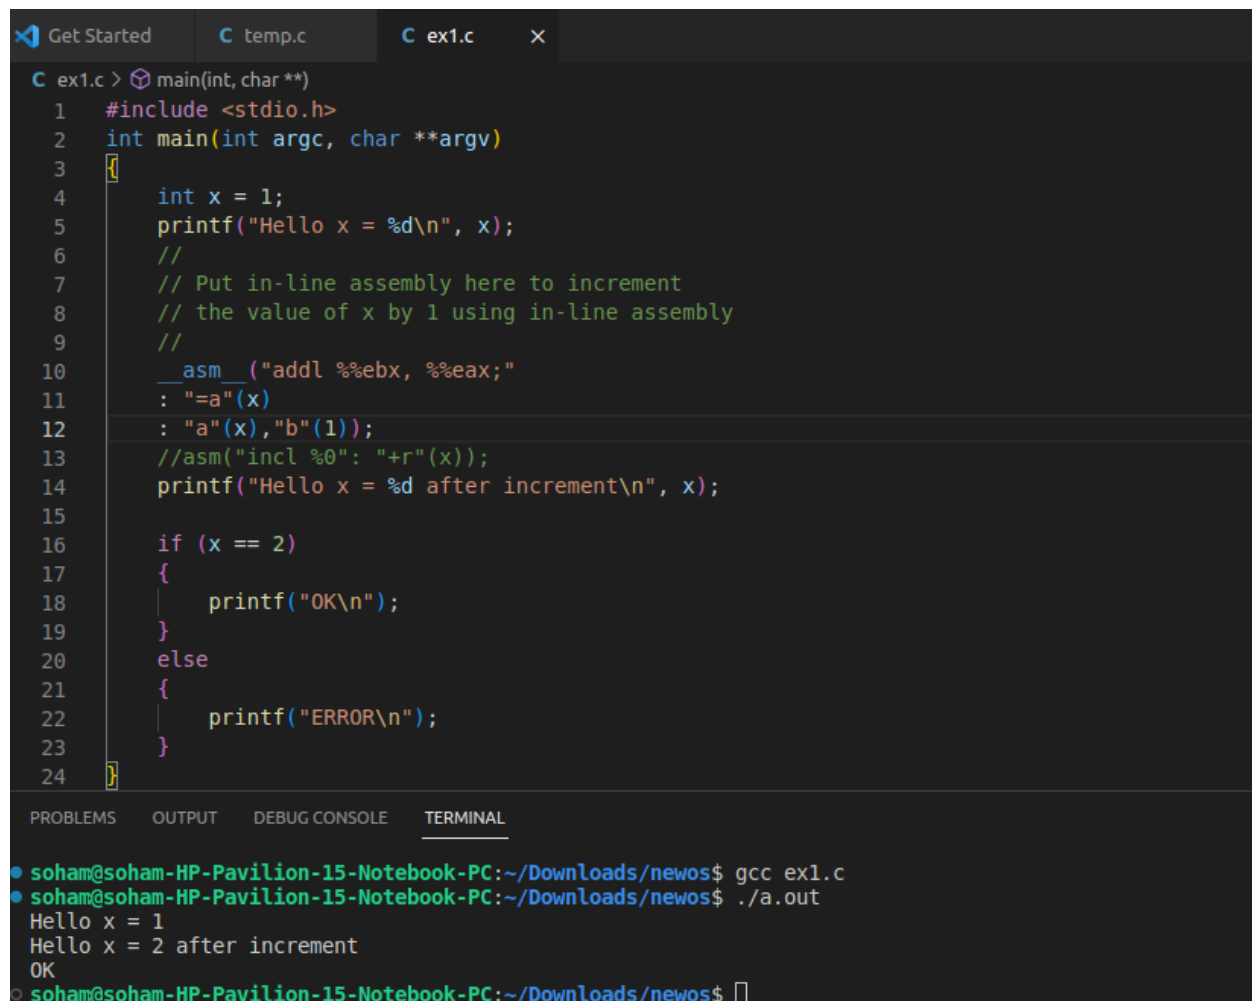


Operating Systems, Lab Report 0A

Soham Roy, Mathematics and Computing
Roll No. 200123055

Exercise 1. Become familiar with inline assembly by writing a simple program. Modify the program ex1.c (at end of this file) to include inline assembly that increments the value of x by 1.

Ex1.c complete code and output



```
C ex1.c > main(int, char **)
1  #include <stdio.h>
2  int main(int argc, char **argv)
3  {
4      int x = 1;
5      printf("Hello x = %d\n", x);
6      //
7      // Put in-line assembly here to increment
8      // the value of x by 1 using in-line assembly
9      //
10     __asm__("addl %%ebx, %%eax;"
11             : "=a"(x)
12             : "a"(x), "b"(1));
13     //asm("incl %0": "+r"(x));
14     printf("Hello x = %d after increment\n", x);
15
16     if (x == 2)
17     {
18         printf("OK\n");
19     }
20     else
21     {
22         printf("ERROR\n");
23     }
24 }
```

PROBLEMS OUTPUT DEBUG CONSOLE TERMINAL

```
soham@soham-HP-Pavilion-15-Notebook-PC:~/Downloads/newos$ gcc ex1.c
soham@soham-HP-Pavilion-15-Notebook-PC:~/Downloads/newos$ ./a.out
Hello x = 1
Hello x = 2 after increment
OK
soham@soham-HP-Pavilion-15-Notebook-PC:~/Downloads/newos$
```

Exercise 2. Use GDB's si (Step Instruction) command to trace into the ROM BIOS for a few more instructions, and try to guess what it might be doing.

1st instruction: [f000:fff0] 0xffff0: ljmp \$0x3630,\$0xf000e05b

- Jump to CS = \$0xf000 & IP = 0xe05b
- 0x3630 is jump to this CS (earlier in the BIOS)
- 0xf000e05b is the IP which is different from the lab because it is 32 bits rather than 16 bits and that is all the way into the top of the extended memory location but before the memory mapped PCI device location reserved by the BIOS

2nd Instruction: [f000:e05b] 0xfe05b: cmpw \$0xffc8,%cs:(%esi)

- Compare content at 0xffc8 & with content at code segment offset with value at esi.
- esi:- 32-bit source index register

3rd Instruction: [f000:e062] 0xfe062: jne 0xd241d0b0

- Jump to 0xd241d0b0 if the above comparison does not set ZF

4th instruction: [f000:e066] 0xfe066: xor %edx,%edx

- ZF was set thus jump of previous instruction doesn't occur
- It set edx to zero, edx is 32-bit general-purpose register.

5th instruction: [f000:e068] 0xfe068: mov %edx,%ss

- Move content of stack segment register(ss) to edx

6th instruction: [f000:e06a] 0xfe06a: mov \$0x7000,%sp

- Move content at the location pointed 16-bit stack pointer(sp) to \$0x7000

```

(gdb) add-auto-load-safe-pathcommand
Undefined command: "add-auto-load-safe-pathcommand". Try "help".
(gdb) si
[f000:e05b] 0xfe05b: cmpw $0xffc8,%cs:(%esi)
0x0000e05b in ?? ()
(gdb) si
[f000:e062] 0xfe062: jne 0xd241d0b0
0x0000e062 in ?? ()
(gdb) si
[f000:e066] 0xfe066: xor %edx,%edx
0x0000e066 in ?? ()
(gdb) si
[f000:e068] 0xfe068: mov %edx,%ss
0x0000e068 in ?? ()
(gdb) si
[f000:e06a] 0xfe06a: mov $0x7000,%sp
0x0000e06a in ?? ()
(gdb) si
[f000:e070] 0xfe070: mov $0xfc1c,%dx
0x0000e070 in ?? ()
(gdb) si
[f000:e076] 0xfe076: jmp 0x5576cf2d

```

Exercise 3:

The code for readsect() is given below :

```

59 void
60 readsect(void *dst, uint offset)
61 {
62     // Issue command.
63     waitdisk();
64     outb(0x1F2, 1); // count = 1
65     outb(0x1F3, offset);
66     outb(0x1F4, offset >> 8);
67     outb(0x1F5, offset >> 16);
68     outb(0x1F6, (offset >> 24) | 0xE0);
69     outb(0x1F7, 0x20); // cmd 0x20 - read sectors
70
71     // Read data.
72     waitdisk();
73     insl(0x1F0, dst, SECTSIZE/4);
74 }

```

The assembly code for readsect() is given below:

```

162 00007c8c <readsect>:
163
164 // Read a single sector at offset into dst.
165 void
166 readsect(void *dst, uint offset)
167 {
168     7c8c:      55                push    %ebp
169     7c8d:      89 e5          mov     %esp,%ebp
170     7c8f:      57                push    %edi
171     7c90:      53                push    %ebx
172     7c91:      8b 5d 0c          mov     0xc(%ebp),%ebx
173     // Issue command.
174     waitdisk();
175     7c94:      e8 e5 ff ff ff    call    7c7e <waitdisk>
176 }
177

```

```

178 static inline void
179 outb(ushort port, uchar data)
180 {
181     asm volatile("out %0,%1" : : "a" (data), "d" (port));
182     7c99:      b8 01 00 00 00      mov     $0x1,%eax
183     7c9e:      ba f2 01 00 00      mov     $0x1f2,%edx
184     7ca3:      ee                      out     %al,(%dx)
185     7ca4:      ba f3 01 00 00      mov     $0x1f3,%edx
186     7ca9:      89 d8                      mov     %ebx,%eax
187     7cab:      ee                      out     %al,(%dx)
188     outb(0x1F2, 1);    // count = 1
189     outb(0x1F3, offset);
190     outb(0x1F4, offset >> 8);
191     7cac:      89 d8                      mov     %ebx,%eax
192     7cae:      c1 e8 08                shr     $0x8,%eax
193     7cb1:      ba f4 01 00 00      mov     $0x1f4,%edx
194     7cb6:      ee                      out     %al,(%dx)
195     outb(0x1F5, offset >> 16);
196     7cb7:      89 d8                      mov     %ebx,%eax
197     7cb9:      c1 e8 10                shr     $0x10,%eax
198     7cbc:      ba f5 01 00 00      mov     $0x1f5,%edx
199     7cc1:      ee                      out     %al,(%dx)
200     outb(0x1F6, (offset >> 24) | 0xE0);
201     7cc2:      89 d8                      mov     %ebx,%eax
202     7cc4:      c1 e8 18                shr     $0x18,%eax
203     7cc7:      83 c8 e0                or      $0xfffffffffe0,%eax
204     7cca:      ba f6 01 00 00      mov     $0x1f6,%edx
205     7ccf:      ee                      out     %al,(%dx)
206     7cd0:      b8 20 00 00 00      mov     $0x20,%eax
207     7cd5:      ba f7 01 00 00      mov     $0x1f7,%edx
208     7cda:      ee                      out     %al,(%dx)
209     outb(0x1F7, 0x20); // cmd 0x20 - read sectors

```

```

207     7cd5:      ba f7 01 00 00      mov     $0x1f7,%edx
208     7cda:      ee                      out     %al,(%dx)
209     outb(0x1F7, 0x20); // cmd 0x20 - read sectors
210
211     // Read data.
212     waitdisk();
213     7cdb:      e8 9e ff ff ff      call    7c7e <waitdisk>
214     asm volatile("cld; rep insl" :
215     7ce0:      8b 7d 08                mov     0x8(%ebp),%edi
216     7ce3:      b9 80 00 00 00      mov     $0x80,%ecx
217     7ce8:      ba f0 01 00 00      mov     $0x1f0,%edx
218     7ced:      fc                      cld
219     7cee:      f3 6d                rep insl (%dx),%es:(%edi)
220     insl(0x1F0, dst, SECTSIZE/4);
221 }

```

The for loop that reads the sectors of the kernel from the disk is given below :

```

37  for(; ph < eph; ph++){
38      pa = (uchar*)ph->paddr;
39      readseg(pa, ph->filesz, ph->off);
40      if(ph->memsz > ph->filesz)
41          stosb(pa + ph->filesz, 0, ph->memsz - ph->filesz);
42  }
43

```

The first instruction of the loop:

```

310      7d7d:      39 f3          cmp     %esi,%ebx

```

The last instruction of the loop:

```

324      7d94:      76 eb          jbe     7d81 <bootmain+0x44>

```

The explanation for the first instruction is that the first operation on entering the for loop will be comparison between the values of ph and eph because the loop will run only when $ph < eph$. The explanation of last instruction is that the loop ends when the values of ph and eph become equal and hence the loop jumps to the next instruction at 0x7d91. Hence the jump instruction will be the last instruction of the for loop. The next instruction after the for loop is

```

313      7d81:      ff 15 18 00 01 00    call    *0x10018

```

Making a breakpoint at that address and then stepping into further instructions gives the following output.

```

(gdb) b *0x7d81
Breakpoint 1 at 0x7d81
(gdb) c
Continuing.
The target architecture is set to "i386".
=> 0x7d81:      call    *0x10018

Thread 1 hit Breakpoint 1, 0x00007d81 in ?? ()
(gdb) si
=> 0x10000c:    mov     %cr4,%eax
0x0010000c in ?? ()
(gdb) si
=> 0x10000f:    or      $0x10,%eax
0x0010000f in ?? ()
(gdb) si
=> 0x100012:    mov     %eax,%cr4
0x00100012 in ?? ()
(gdb) si
=> 0x100015:    mov     $0x10a000,%eax
0x00100015 in ?? ()
(gdb) si
=> 0x10001a:    mov     %eax,%cr3
0x0010001a in ?? ()
(gdb) si
=> 0x10001d:    mov     %cr0,%eax
0x0010001d in ?? ()
(gdb) si
=> 0x100020:    or      $0x80010000,%eax
0x00100020 in ?? ()
(gdb) si
=> 0x100025:    mov     %eax,%cr0
0x00100025 in ?? ()

```

(a)

The command line `$(SEG_KCODE<<3), $start32` causes the switch from 16 to 32-bit mode in `bootasm.S`

```
0x00100028 in ?? ()
(gdb) si
=> 0x100028:    mov     $0x801164d0,%esp
0x00100028 in ?? ()
(gdb)
=> 0x10002d:    mov     $0x80103060,%eax
0x0010002d in ?? ()
(gdb) si
=> 0x100032:    jmp     *%eax
0x00100032 in ?? ()
(gdb) si
=> 0x80103060 <main>:  lea     0x4(%esp),%ecx
main () at main.c:20
20      kinit1(end, P2V(4*1024*1024)); // phys page allocator
(gdb) █
```



```
soham@soham-HP-Pavilion-15-Notebook-PC: ~/xv6-public
of GDB. Attempting to continue with the default i8086 settings.
(gdb) si
[f000:e05b] 0xfe05b: cmpw $0xffc8,%cs:(%esi)
0x0000e05b in ?? ()
(gdb) b *0x7c00
Breakpoint 1 at 0x7c00
(gdb) c
Continuing.
[ 0:7c00] => 0x7c00: cli

Thread 1 hit Breakpoint 1, 0x00007c00 in ?? ()
(gdb) si
[ 0:7c01] => 0x7c01: xor %eax,%eax
0x00007c01 in ?? ()
(gdb) si
[ 0:7c03] => 0x7c03: mov %eax,%ds
0x00007c03 in ?? ()
(gdb)
[ 0:7c05] => 0x7c05: mov %eax,%es
0x00007c05 in ?? ()
(gdb) x/20i 0x7c00
0x7c00: cli
0x7c01: xor %eax,%eax
0x7c03: mov %eax,%ds
=> 0x7c05: mov %eax,%es
0x7c07: mov %eax,%ss
0x7c09: in $0x64,%al
0x7c0b: test $0x2,%al
0x7c0d: jne 0x7c09
0x7c0f: mov $0xd1,%al
0x7c11: out %al,$0x64
0x7c13: in $0x64,%al
0x7c15: test $0x2,%al
0x7c17: jne 0x7c13
0x7c19: mov $0xdf,%al
0x7c1b: out %al,$0x60
0x7c1d: lgdtl (%esi)
0x7c20: js 0x7c9e
0x7c22: mov %cr0,%eax
0x7c25: or $0x1,%ax
0x7c29: mov %eax,%cr0
(gdb)
```

(b)

By analysing the contents of bootasm.S, bootmain.c and bootblock.asm, we conclude that bootasm.S switches the OS into 32-bit mode and then calls bootmain.c which first loads the kernel using ELF header and then enters the kernel using entry(). Hence the last instruction of bootloader is entry(). Looking for the same in bootblock.asm, we find out the instruction to be

```

312  entry();
313  7d81:      ff 15 18 00 01 00      call  *0x10018

```

which is a call instruction which shifts control to the address stored at 0x10018 since dereferencing operator (*) has been used. Now we need to know the starting address of the kernel. We can find this by two methods:

- (i) By looking at the first word of memory stored at 0x10018 (by using the command “x/1x 0x10018”)
- (ii) By looking at the contents of “objdump -f kernel”

After getting the starting address of kernel, we need to see what is the instruction stored at that address to get the first instruction of kernel. We can do this by two methods:

- (i) By using “x/1i 0x0010000c”
- (ii) By looking into kernel.asm

```

soham@soham-HP-Pavilion-15-Notebook-PC: ~/xv6-public
(gdb) b *0x7d81
Breakpoint 3 at 0x7d81
(gdb) c
Continuing.
=> 0x7d81:      call  *0x10018

Thread 1 hit Breakpoint 3, 0x00007d81 in ?? ()
(gdb) x/1x 0x10018
0x10018:      0x0010000c
(gdb) x/1i 0x0010000c
0x10000c:      mov    %cr4,%eax
(gdb)

```

Hence, the first instruction of kernel is

```

0x10000c:      mov    %cr4,%eax

```

The above lines of code are present in bootmain.c. This is the code that is used by xv6 to load the kernel. xv6 first loads ELF headers of kernel into a memory location pointed to by “elf”. Then it stores the starting address of the first segment of the kernel to be loaded in “ph” by adding an offset (“elf->phoff”) to the starting address (elf). It also maintains an end pointer eph which points to the memory location after the end of the last segment. It then iterates over all the segments. For every segment, pa points to the address at which this segment has to be loaded. Then it loads the current segment at that location by passing pa, ph->filesz and ph->off

parameters to readseg. It then checks the memory assigned to this sector is greater than the data copied. If this is true, it initializes the extra memory with zeros.

Coming back to the question, the boot loader keeps loading segments while the condition “ph < eph” is true. The values of ph and eph are determined using attributes phoff and phnum of the ELF header. So, the information stores in the ELF header helps the boot loader to decide how many sectors it has to read.

Exercise 4. Read about programming with pointers in C. Then download the code for pointers.c, run it, and make sure you understand where all of the printed values come from. In particular, make sure you understand where the pointer addresses in lines 1 and 6 come from, how all the values in lines 2 through 4 get there, and why the values printed in line 5 are seemingly corrupted. We also recommend reading the K-splice pointer challenge as a way to test that you understand how pointer arithmetic and arrays work in C.

Answer:

Output of code in pointer.c

```

#include <stdio.h>
#include <stdlib.h>

void f(void){
    int a[4];
    int *b=malloc(16);
    int *c;
    int i;
    printf("1: a=%p, b=%p,c=%p\n",a,b,c);
    c=a;
    for(int i=0;i<4;i++){
        a[i]=100+i;
    }
    c[0]=200;
    printf("2: a[0]=%d,a[1]= %d, a[2]=%d,a[3]= %d\n",a[0],a[1],a[2],a[3]);
    c[1]=300;
    *(c+2)=301;
    3[c]=302;
    printf("3: a[0] =%d,a[1]= %d,a[2]=%d,a[3]=%d\n",a[0],a[1],a[2],a[3]);
    c=c+1;
    *c=400;
    printf("4: a[0] =%d,a[1]= %d,a[2]=%d,a[3]=%d\n",a[0],a[1],a[2],a[3]);
    c=(int *)((char *)c+1);
    *c=500;
    printf("5: a[0] =%d,a[1]= %d,a[2]=%d,a[3]=%d\n",a[0],a[1],a[2],a[3]);
    b=(int *)a+1;
    c=(int *)((char*)a+1);
    printf("6: a=%p, b=%p, c=%p\n",a,b,c);
}

int main(int ac,char **av){
    f();
    return 0;
}

```

```

● soham@soham-HP-Pavilion-15-Notebook-PC:~/Downloads/newos$ gcc temp.c
● soham@soham-HP-Pavilion-15-Notebook-PC:~/Downloads/newos$ ./a.out
1: a=0x7fffd689d7a0, b=0x5562c19c42a0,c=(nil)
2: a[0]=200,a[1]= 101, a[2]=102,a[3]= 103
3: a[0] =200,a[1]= 300,a[2]=301,a[3]=302
4: a[0] =200,a[1]= 400,a[2]=301,a[3]=302
5: a[0] =200,a[1]= 128144,a[2]=256,a[3]=302
6: a=0x7fffd689d7a0, b=0x7fffd689d7a4, c=0x7fffd689d7a1
○ soham@soham-HP-Pavilion-15-Notebook-PC:~/Downloads/newos$ █

```

objdump -h kernel

```
soham@soham-HP-Pavilion-15-Notebook-PC:~/xv6-public$ objdump -h kernel
```

```
kernel:      file format elf32-i386
```

```
Sections:
```

Idx	Name	Size	VMA	LMA	File off	Algn
0	.text	00007188	80100000	00100000	00001000	2**4
	CONTENTS, ALLOC, LOAD, READONLY, CODE					
1	.rodata	000009cb	801071a0	001071a0	000081a0	2**5
	CONTENTS, ALLOC, LOAD, READONLY, DATA					
2	.data	00002516	80108000	00108000	00009000	2**12
	CONTENTS, ALLOC, LOAD, DATA					
3	.bss	0000afb0	8010a520	0010a520	0000b516	2**5
	ALLOC					
4	.debug_line	00006aaf	00000000	00000000	0000b516	2**0
	CONTENTS, READONLY, DEBUGGING, OCTETS					
5	.debug_info	00010e14	00000000	00000000	00011fc5	2**0
	CONTENTS, READONLY, DEBUGGING, OCTETS					
6	.debug_abbrev	00004496	00000000	00000000	00022dd9	2**0
	CONTENTS, READONLY, DEBUGGING, OCTETS					
7	.debug_aranges	000003b0	00000000	00000000	00027270	2**3
	CONTENTS, READONLY, DEBUGGING, OCTETS					
8	.debug_str	00000dee	00000000	00000000	00027620	2**0
	CONTENTS, READONLY, DEBUGGING, OCTETS					
9	.debug_loclists	000050b1	00000000	00000000	0002840e	2**0
	CONTENTS, READONLY, DEBUGGING, OCTETS					
10	.debug_rnglists	00000845	00000000	00000000	0002d4bf	2**0
	CONTENTS, READONLY, DEBUGGING, OCTETS					
11	.debug_line_str	00000131	00000000	00000000	0002dd04	2**0
	CONTENTS, READONLY, DEBUGGING, OCTETS					
12	.comment	00000026	00000000	00000000	0002de35	2**0
	CONTENTS, READONLY					

As we can see in the above screenshot, VMA and LMA of .text section is different indicating that it loads and executes from different addresses. `objdump -h bootblock.o`

```

bootblock.o:      file format elf32-l386

Sections:
Idx Name          Size      VMA       LMA       File off  Algn
  0 .text          000001c3  00007c00  00007c00  00000074  2**2
    CONTENTS, ALLOC, LOAD, CODE
  1 .eh_frame      000000b0  00007dc4  00007dc4  00000238  2**2
    CONTENTS, ALLOC, LOAD, READONLY, DATA
  2 .comment       00000026  00000000  00000000  000002e8  2**0
    CONTENTS, READONLY
  3 .debug_aranges 00000040  00000000  00000000  00000310  2**3
    CONTENTS, READONLY, DEBUGGING, OCTETS
  4 .debug_info    00000585  00000000  00000000  00000350  2**0
    CONTENTS, READONLY, DEBUGGING, OCTETS
  5 .debug_abbrev  0000023c  00000000  00000000  000008d5  2**0
    CONTENTS, READONLY, DEBUGGING, OCTETS
  6 .debug_line    00000283  00000000  00000000  00000b11  2**0
    CONTENTS, READONLY, DEBUGGING, OCTETS
  7 .debug_str     0000020b  00000000  00000000  00000d94  2**0
    CONTENTS, READONLY, DEBUGGING, OCTETS
  8 .debug_line_str 00000046  00000000  00000000  00000f9f  2**0
    CONTENTS, READONLY, DEBUGGING, OCTETS
  9 .debug_loclists 0000018d  00000000  00000000  00000fe5  2**0
    CONTENTS, READONLY, DEBUGGING, OCTETS
 10 .debug_rnglists 00000033  00000000  00000000  00001172  2**0
    CONTENTS, READONLY, DEBUGGING, OCTETS

```

As we can see in the above screenshot, VMA and LMA of .text section is same indicating that it loads and executes from the same address.

Exercise 5. Trace through the first few instructions of the boot loader again and identify the first instruction that would "break" or otherwise do the wrong thing if you were to get the boot loader's link address wrong. Then change the link address in Makefile to something wrong, run make clean, recompile the lab with make, and trace into the boot loader again to see what happens. Don't forget to change the link address back and make clean again afterwards!

Look back at the load and link addresses for the kernel. Unlike the boot loader, these two addresses aren't the same: the kernel is telling the boot loader to load it into memory at a low address (1 MB), but it expects to execute from a high address. We'll dig in to how we make this work in the next section.

Besides the section information, there is one more field in the ELF header that is important to us, named `e_entry`. This field holds the link address of the entry

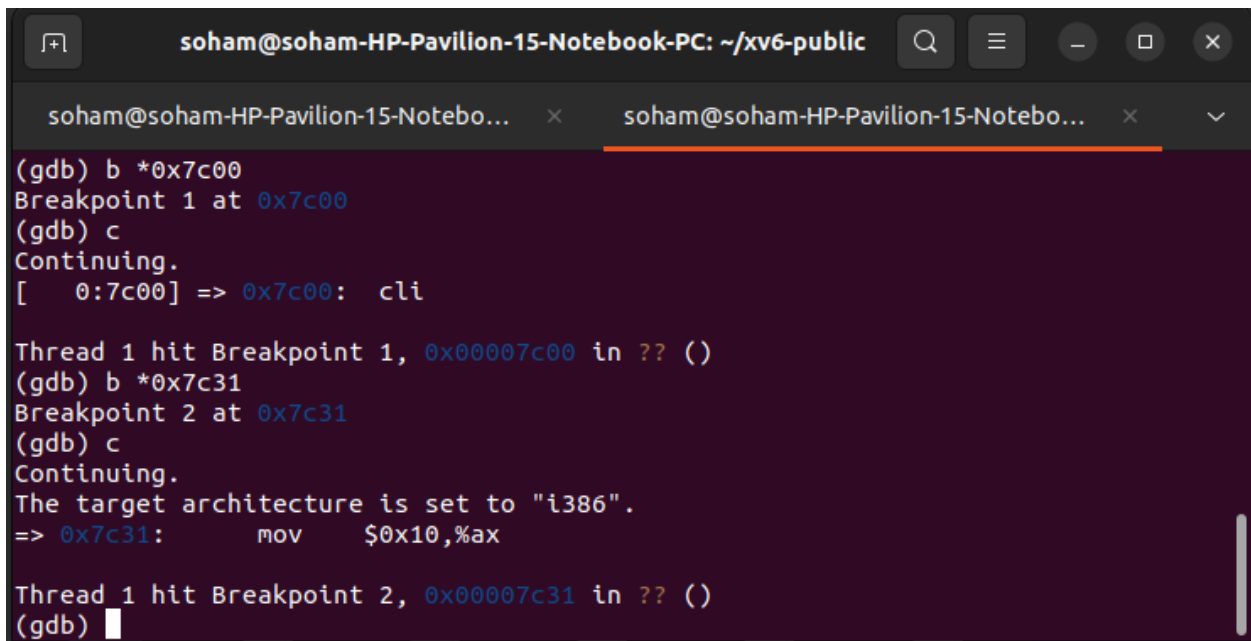
point in the program: the memory address in the program's text section at which the program should begin executing. You can see the entry point:

\$ objdump -f kernel

You should now be able to understand the minimal ELF loader in bootmain.c. It reads each section of the kernel from disk into memory at the section's load address and then jumps to the kernel's entry point.

Answer:

When boot loader's link address is **0x7C00** then commands are running properly and transition from 16 to 32 bit was occurring at **0x7C31** address location as seen below:



```
(gdb) b *0x7c00
Breakpoint 1 at 0x7c00
(gdb) c
Continuing.
[ 0:7c00] => 0x7c00: cli

Thread 1 hit Breakpoint 1, 0x00007c00 in ?? ()
(gdb) b *0x7c31
Breakpoint 2 at 0x7c31
(gdb) c
Continuing.
The target architecture is set to "i386".
=> 0x7c31: mov $0x10,%ax

Thread 1 hit Breakpoint 2, 0x00007c31 in ?? ()
(gdb)
```

But when the boot loader's link address is changed to any other address (I took **0x7C24** in this case), after running

make clean

make

and restarting gdb

and continuing from address location 0x7C00,

then the boot loader is restarting again and again after running some instructions in the gdb.

```
soham@soham-HP-Pavilion-15-Notebo... x soham@soham-HP-Pavilion-15-Notebo... x v
(gdb) b *0x7c00
Breakpoint 1 at 0x7c00
(gdb) c
Continuing.
[ 0:7c00] => 0x7c00: cli

Thread 1 hit Breakpoint 1, 0x00007c00 in ?? ()
(gdb) c
Continuing.
[ 0:7c00] => 0x7c00: cli

Thread 1 hit Breakpoint 1, 0x00007c00 in ?? ()
(gdb) b *0x7c55
Breakpoint 2 at 0x7c55
(gdb) c
Continuing.
[ 0:7c00] => 0x7c00: cli

Thread 1 hit Breakpoint 1, 0x00007c00 in ?? ()
(gdb) si
[ 0:7c01] => 0x7c01: xor    %eax,%eax
0x00007c01 in ?? ()
(gdb) si
[ 0:7c03] => 0x7c03: mov    %eax,%ds
0x00007c03 in ?? ()
(gdb) c
Continuing.
[ 0:7c00] => 0x7c00: cli

Thread 1 hit Breakpoint 1, 0x00007c00 in ?? ()
(gdb)
```

As seen in the image above, we tried to run commands after continuing from breakpoint at **0x7C00** address location and we always end up hitting the same breakpoint at **0x7C00**. Also 16-to-32-bit architecture change didn't occur as breakpoint **b *0x7C55** is not hit which should be responsible for architecture change in this case.

ljmp \$(SEG_KCODE<<3), \$start32 is the first instruction that breaks. Before changing the link address of the boot loader, from address **0x7C00**, after performing 2-3 si 10 instructions, architecture changed from 16 to 32 bit.

But after changing the link address to **0x7C24**, architecture didn't change which means that the boot loader is not loaded properly at the changed link address.

Exercise 6:

Answer:

At the point when BIOS enters the boot loader (at first breakpoint):

```
soham@soham-HP-Pavilion-15-Notebook-PC: ~/xv6-public
soham@soham-HP-Pavilion-15-Notebo... x soham@soham-HP-Pavilion-15-Notebo... x v
(gdb) si
[f000:e05b] 0xfe05b: cmpw $0xffc8,%cs:(%esi)
0x0000e05b in ?? ()
(gdb) b *0x7c00
Breakpoint 1 at 0x7c00
(gdb) c
Continuing.
[ 0:7c00] => 0x7c00: cli

Thread 1 hit Breakpoint 1, 0x00007c00 in ?? ()
(gdb) x/8x 0x00100000
0x100000: 0x00000000 0x00000000 0x00000000 0x00000000
0x100010: 0x00000000 0x00000000 0x00000000 0x00000000
(gdb) x/8i 0x00100000
0x100000: add %al,(%eax)
0x100002: add %al,(%eax)
0x100004: add %al,(%eax)
0x100006: add %al,(%eax)
0x100008: add %al,(%eax)
0x10000a: add %al,(%eax)
0x10000c: add %al,(%eax)
0x10000e: add %al,(%eax)
(gdb)
```

At the point when the boot loader enters the kernel (at second breakpoint):

```
soham@soham-HP-Pavilion-15-Notebook-PC: ~/xv6-public
soham@soham-HP-Pavilion-15-Notebo... x soham@soham-HP-Pavilion-15-Notebo... x
0x10000a: add    %al, (%eax)
0x10000c: add    %al, (%eax)
0x10000e: add    %al, (%eax)
(gdb) b *0x0010000c
Breakpoint 2 at 0x10000c
(gdb) c
Continuing.
The target architecture is set to "i386".
=> 0x10000c: mov    %cr4, %eax

Thread 1 hit Breakpoint 2, 0x0010000c in ?? ()
(gdb) x/8x 0x00100000
0x100000: 0x1badb002    0x00000000    0xe4524ffe    0x83e0200f
0x100010: 0x220f10c8    0xa000b8e0    0x220f0010    0xc0200fd8
(gdb) x/8i 0x00100000
0x100000: add    0x1bad(%eax), %dh
0x100006: add    %al, (%eax)
0x100008: decb   0x52(%edi)
0x10000b: in     $0xf, %al
0x10000d: and    %ah, %al
0x10000f: or     $0x10, %eax
0x100012: mov    %eax, %cr4
0x100015: mov    $0x10a000, %eax
(gdb)
```

8 words of instruction at 0x00100000 at the point when BIOS enters the boot loader and 8 words of instruction at 0x00100000 at the point when the boot loader enters the kernel are different as when the BIOS enters and loads the boot loader, then it just loads it in memory location between 0x7C00 and 0x7DFF due to which all the 8 words of instructions are zero at 0x00100000. But before the boot loader enters the kernel, it already has performed the 16-to-32-bit transition and setting up of stack which leads to new instructions at address 0x00100000.